

# Project Report – Group 3

Jasper Bruin, Richard Wijers, Alexandru Dumitriu

## 1 INTRODUCTION

This is a minimal version of ACM’s consolidated article template that can be used for reports in classes. Styles are original, but some information, like the conference, keywords, or other overhead that is only relevant for publishing, has been removed through redefinition of commands in the beginning of the document.

This document can be used like any other latex document. Please note that you *MUST NOT* change the formatting (e.g., font sizes, line heights, margins, ...). However, the *structure* of the document (e.g., abstract, introduction, ...) usually needs editing to reflect the assignment instructions.

## 2 RELEASE PIPELINE DOCUMENTATION

### 2.1 Software Package Release Pipeline

**2.1.1 Overview.** This section documents the release pipeline for the Python package `lib-ml`. It describes each step, the tools involved, and the data flow throughout the process. The pipeline is activated by pushes to the main branch, creation of tags prefixed with ‘v’, and pull requests targeting the main branch. This configuration ensures that new versions undergo a rigorous build and test process before publication.

#### 2.1.2 Pipeline Steps.

- (1) Checkout Repository
- (2) Set up Python Environment
- (3) Install Poetry
- (4) Configure Poetry
- (5) Install Dependencies
- (6) Code Analysis with Pylint
- (7) Code Analysis with Flake8
- (8) Check Version on PyPI
- (9) Bump Version and Tag
- (10) Build and Publish to PyPI
- (11) Clean Up

**2.1.3 Purpose and Implementation.** The pipeline begins with checking out the repository using `actions/checkout@v4`. The purpose is to fetch the latest code from the repository, and it is implemented by cloning the repository into the runner environment with the full history by setting `fetch-depth` to 0. Next, the Python environment is set up using `actions/setup-python@v3` with Python version 3.9. This step is crucial for ensuring the correct Python version is used throughout the pipeline.

The Poetry package manager is then installed to handle dependencies and packaging. This is implemented by running a command to download and install Poetry. Following the installation, Poetry is configured to not create virtual environments, implemented by running `poetry config virtualenvs.create false`, which simplifies the environment setup.

Dependencies specified in `pyproject.toml` are installed using `poetry install`. This step is essential for fetching and installing all required dependencies for the project. The code is then analyzed

for errors and compliance with coding standards using `pylint` and `flake8`. These steps are implemented by running `pylint` and `flake8` commands on the Python files to ensure code quality and style adherence.

A crucial step involves checking if the current version of the package already exists on PyPI. This is done by querying PyPI with the current version obtained from Poetry, and an output variable is set based on the version’s existence. If the version does not exist on PyPI, the project version is incremented, and the repository is tagged with the new version. This involves updating `pyproject.toml`, committing the changes, tagging the new version, and pushing these updates to the repository, which is implemented through a series of git commands.

The next step is to build the package and publish it to PyPI. This is accomplished by running `poetry build` to create distribution artifacts (`sdist` and `wheel`) and then publishing them using provided credentials. Finally, the pipeline performs a clean-up by removing the `dist` directory to maintain a clean working environment.

**2.1.4 Data Flow.** The data flows from the repository being cloned into the CI environment, through setting up the Python environment and installing dependencies, to analyse the code. If the version is new, the versioning data is updated, and the code changes are pushed back to the repository. The final artifact, a Python package, is built and published to PyPI. The clean-up step ensures no residual data is left in the environment.

**2.1.5 Tools Used.** The tools integral to this pipeline include GitHub Actions, Python, and Poetry for Python dependency management. Code analysis is performed using Pylint and Flake8.

#### 2.1.6 Artifacts.

- Python package distribution files (`.tar.gz`, `.whl`)
- Published package on PyPI

### 2.2 Container Image Release Pipeline

**2.2.1 Overview.** The pipeline is triggered by pushes to the main branch or the creation of tags starting with ‘v’. This ensures that the Docker image is updated and deployed in response to changes in the source code or new version releases.

#### 2.2.2 Pipeline Steps.

- (1) Checkout Repository
- (2) Build Docker Image
- (3) Log in to Registry
- (4) Push Docker Image

**2.2.3 Purpose and Implementation.** The Docker Publish pipeline starts by checking out the repository using `actions/checkout@v4`. The purpose is to fetch the latest code from the repository, implemented by cloning the repository into the runner environment. Following this, the Docker image is built from the Dockerfile. This is implemented by running the command `docker build . -file`

`Dockerfile -tag \${IMAGE_NAME}`, which creates a Docker image tagged with `\${IMAGE_NAME}`.

Once the image is built, the pipeline logs in to the Docker registry using the GitHub token to authenticate. This step is implemented by echoing the GitHub token and passing it to `docker login`, establishing a session with the Docker registry. The final step involves tagging the built Docker image with the appropriate version and pushing it to the registry. The version is derived from the git ref, and if the ref corresponds to a tag, the version is adjusted accordingly. The image is then pushed to the registry under the derived tag, implemented by using Docker commands to tag and push the image.

**2.2.4 Data Flow.** The data flow in this pipeline starts with the GitHub repository as the source, moves through the GitHub Actions runner environment where the Docker image is built and ends at the GitHub Container Registry where the image is stored. This streamlined process ensures that new or updated Docker images are readily available for deployment or further development stages.

**2.2.5 Tools Used.** Throughout this process, tools such as Docker for building and pushing images, and GitHub Actions for automating these steps, play a crucial role.

**2.2.6 Artifacts:**

- Docker image - ghcr.io

### 3 DEPLOYMENT DOCUMENTATION

- Visual representation of the deployment structure
- Description of data flow for incoming requests
- Overview of all resources and their connections

### 4 EXTENSION PROPOSAL

- Identification of project shortcomings
- Proposed extensions or refactoring
- Sources and inspirations for the solution
- Testing of the new design

### 5 ISTIO USE CASE

- General description of the use case
- Changes from the base design

### 6 ISTIO EXPERIMENTAL SETUP

- Description of the experiment
- Hypothesis and metrics
- Data decision process

## 7 ML PIPELINE DOCUMENTATION

This section provides a overview of the Machine Learning (ML) pipeline used in our project. It covers the setup and configuration details, the decision-making process and design choices, automated tasks, and created artefacts.

### 7.1 Tools used

The project uses the Cookiecutter project template. Cookiecutter standardizes the project setup process with predefined directory

structures and configuration files. This helps to maintain consistency and predictability for the project structure.

For managing Python dependencies, we use Poetry. Poetry simplifies dependency management by ensuring that all collaborators use the same versions of dependencies across different environments. It also automates the creation and management of virtual environments, which helps to isolate the project and its dependencies from other Python projects on the same machine.

To ensure code quality and consistency, we use Pylint and Flake8 as linters. Pylint checks for errors, bugs, and code that does not adhere to best practices, while Flake8 enforces style guidelines. This helps us maintain a high standard of code quality and readability.

Data Version Control (DVC) is an open-source tool that is designed to handle large data files, binary models, and metrics as well as code. It is built to make Machine Learning (ML) projects reproducible and shareable. DVC is compatible with Git repositories and has a similar command-line interface.

### 7.2 Pipeline stages

The ML pipeline consists of four main stages as shown in figure 1. each interacting or producing some result as follows:

- (1) **download\_data:** This stage is responsible for downloading the dataset. It depends on the `download_dataset.py` script and produces the output in the `data/external` directory. The parameters `data_folder_id` and `data_folder` are used in this stage.
- (2) **preprocess\_data:** This stage is responsible for tokenizing the data. It depends on the `pyproject.toml` file and produces the output in the `artifacts/tokenized` directory. The parameters `tokenized_folder` and `data_folder` are used in this stage.
- (3) **train\_model:** This stage is responsible for training the model. It depends on the `train_model.py` script and produces the output in the `artifacts/trained` directory. The parameters `epochs`, `batch_size`, `categories`, `loss_function`, `optimizer`, `tokenized_folder`, and `trained_folder` are used in this stage.
- (4) **predict\_model:** This stage is responsible for making predictions using the trained model. It depends on the `predict_model.py` script and produces the output in the `artifacts/predicted` directory. The parameters `predicted_folder`, `trained_folder`, and `tokenized_folder` are used in this stage.

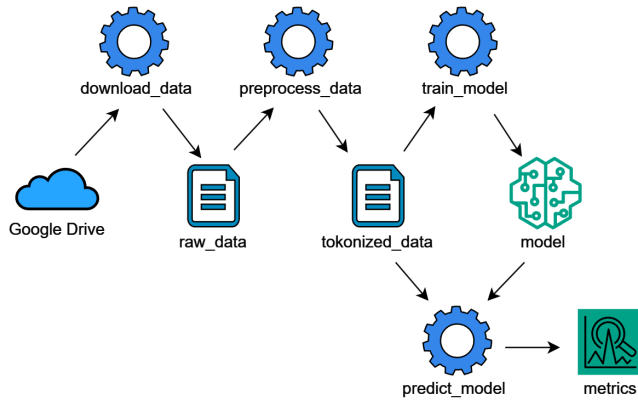


Figure 1: ML Pipeline

### 7.3 Automation

To ensure quality, two automatic GitHub actions are used:

- `pylint.yml` defines a GitHub Actions workflow that checks Python code quality and linting on every push event. It runs on the latest Ubuntu, sets up Python 3.9, installs and configures Poetry for dependency management, installs project dependencies, and then runs Pylint and Flake8 for static code analysis. Pylint checks for programming errors, while Flake8 enforces coding style guidelines.
- The `pytest.yml` file defines a GitHub Actions workflow that runs unit tests on every pull request. It operates on the latest Ubuntu, sets up Python 3.9, checks out the code, installs dependencies using Poetry, downloads the dataset by running a Python script, and then runs the unit tests using pytest. This ensures that any changes made in a pull request do not break the functionality of the existing code.

### 7.4 Artifacts

The pipeline produces several artifacts that are used in different stages of the pipeline or for evaluation purposes. These include the downloaded data, tokenized data, trained model, and prediction results. These artifacts are stored in the artifacts directory, with each type of artifact stored in its respective subdirectory (tokenized, trained, predicted). This organization makes it easy to locate and use the artifacts in different stages of the pipeline or for evaluation purposes.

## 8 TESTING STRATEGIES ANALYSIS

Table 1 outlines our comprehensive testing strategy. Each type of test and phase in the table is mentioned for quick reference. The table systematically categorizes the sequence of testing phases: Initiate, Data Validation, Feature Handling, Evaluation, Assessment, Infrastructure, Monitoring, Completion, and End. This structured approach ensures that each test is properly set up, executed, reviewed for performance, and analyzed for outcomes. Detailed information about these tests can be found in our `model-training` repository.

### 8.1 Automation of Testing Processes

Our testing processes are fully automated to ensure reliability and efficiency. We leverage our Data Version Control (DVC) pipeline to manage and download datasets from remote sources. This automation ensures that the most current and relevant data is used during testing, reflecting real-world conditions as closely as possible.

**8.1.1 Continuous Integration Setup.** We employ GitHub Actions to automate our testing workflow, as detailed in our `pytest.yml` configuration. This setup is triggered on every pull request, ensuring that all changes are thoroughly tested before integration. Here's how our testing pipeline operates:

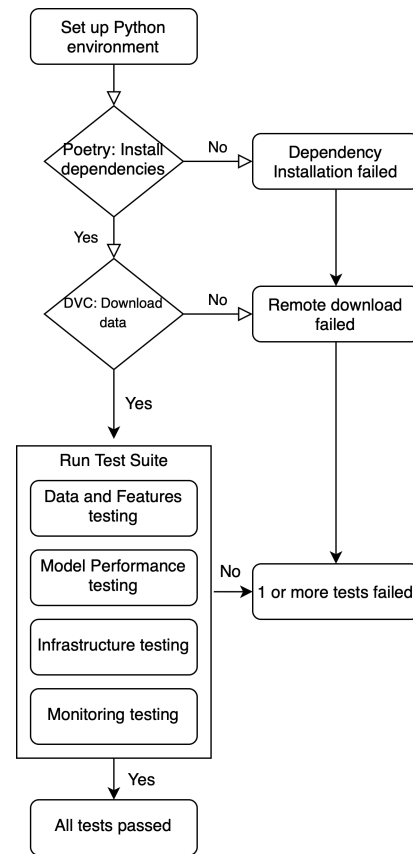


Figure 2: Workflow of the automated CI Testing after each Pull Request

- **Code Checkout:** GitHub Actions checks out the latest code for testing.
- **Environment Setup:** Sets up Python and installs dependencies using Poetry, which ensures that our testing environment is consistent and replicable.
- **Dataset Download:** The dataset is automatically downloaded using a script facilitated by our DVC setup. This step is for testing with the most up-to-date data.
- **Test Execution:** Tests are run using Pytest, which systematically executes various test cases, including unit tests, integration tests, and system tests.

Phase	Tasks
Data Validation	Ensure Data Integrity, Efficient Data Reading, Manage File Operations
Feature Handling	Incorporate Randomness, Remove Sensitive Information, Verify Tokenization and Encoding
Evaluation	Feature Tests, Evaluate Data Slice Mechanisms
Assessment	Check Feature Distribution, Examine Feature-Target Relationships, Assess Computational Cost, Uphold Privacy Considerations, Model Performance Tests
Infrastructure	Verify GPU Resources, Conduct Memory Usage Tests, Validate Model Persistence, Integration Tests, Monitoring Tests
Monitoring	Develop Monitoring Tests, Track Performance Metrics, Detect Anomalies
Completion	Trigger Alerts or Retraining, Results Analysis

Table 1: Workflow Overview

## 8.2 Tests for Features and Data

Our testing framework is designed to ensure data integrity and feature quality, particularly for the detection of phishing URLs. This section explores the specific testing procedures, highlighting their critical roles and outcomes with brief code snippets.

**8.2.1 Data Handling and Random Sampling.** We have developed functions to efficiently read and randomly sample data from datasets. These functions not only streamline the process of loading data but also incorporate randomness to simulate real-world scenarios. For example, data sampling is performed using the `random.shuffle` method, ensuring unbiased data selection:

```
random.shuffle(lines)
```

**8.2.2 Sensitive Information Removal.** Sensitive data is redacted from URLs using a regex pattern defined as follows:

```
re.compile(r"(@|token|session|user|userid|"  
"password|auth|files|pro)")
```

This ensures compliance with data privacy standards and prevents any potential data leakage.

**8.2.3 Tokenization and Label Encoding.** The robustness of our tokenization and label encoding processes is thoroughly tested. The tokenizer transforms URLs into a machine-readable format, which is verified by checking:

```
assert len(data["tokenizer"].word_index) > 0
```

Similarly, label encoding converts categorical labels into numerical format, ensuring the model receives correctly formatted input.

**8.2.4 Feature Distribution and Target Relationship Analysis.** We evaluate the distribution of features and examine the relationships between features and targets. This includes checking that feature lengths are consistent and meaningful:

```
sequence_lengths = [len(seq) for seq in sequences]  
assert np.mean(sequence_lengths) > 0
```

Such tests confirm the logical alignment of data inputs to their respective outputs, essential for effective learning.

**8.2.5 Computational Cost Assessment.** The computational efficiency of processing features is assessed by evaluating the size of the tokenizer's word index:

```
assert len(data["tokenizer"].word_index) < 10000
```

This step ensures that the memory usage and processing time remain within practical limits for real-time applications.

These tests, structured with direct references to the code logic, validate not only the functionality and efficiency of our data processing methods but also uphold stringent privacy standards, ensuring that our phishing detection system is both effective and ethical.

## 8.3 Tests for Model Development

In the development of our URL Publishing Detection CNN Phishing Detector, we implement a series of rigorous tests aimed at evaluating model performance metrics including accuracy, precision, recall, and F1 scores. These metrics are computed for different data subsets—short, medium, and long text slices—to ensure that our model performs consistently across various types of input data.

**8.3.1 Accuracy Testing Across Text Lengths.** To assess the robustness and uniformity of our model's predictive accuracy, we perform tests across three defined text length categories: short, medium, and long. We ensure that the model's accuracy does not vary significantly between these categories, which is crucial for maintaining reliable detection across different URL lengths. The test checks if the differences in accuracy are within acceptable limits:

```
assert abs(short_metrics['accuracy'] -  
medium_metrics['accuracy']) < 0.25
```

```
assert abs(medium_metrics['accuracy'] -  
long_metrics['accuracy']) < 0.25
```

From the results in 3, we can see that short slices exhibit significant variability in all metrics, indicating that the model may struggle with shorter data segments. This could be due to insufficient information within short inputs, leading to inconsistent performance. Medium and long slices demonstrate consistently high performance across all metrics. This suggests that the model performs well when given more context or information within the input data, resulting in better precision, recall, and overall accuracy. The F1 Score for medium and long slices remains high and stable, indicating a well-balanced trade-off between precision and recall, which is crucial for reliable model performance.

**8.3.2 Precision, Recall and F1-Evaluation.** Precision and recall are measured to ensure that the model not only accurately identifies phishing attempts but also minimizes false positives and negatives,

Test Name	Statistic	Expected Value/Condition	Result
Tokenizer	Number of sequences	Match number of input texts	300
	Word index	Non-empty	1659
Label Encoder	Transformed labels	Match number of input labels	300
	Unique classes	More than one	2
Feature Distribution	Mean sequence length	Greater than zero	11.33
Feature-Target Relationship	Number of sequences	Match number of labels	300
Feature Cost	Feature size	Less than 10,000	1874
Feature Privacy	Sensitive patterns	Not present in word index	True
Feature Code	URL cleaning	All URLs start with 'http://' or 'https://'	True
	Input feature type	All features are strings	True
	Input feature length	No empty strings	True

Table 2: Summary of Test Statistics and Expected Conditions

Slice	Accuracy	Precision	Recall	F1 Score
Short	0.88	0.8864	0.88	0.8817
Short	0.6957	0.7846	0.6957	0.7075
Short	0.7647	0.8371	0.7647	0.7469
Short	0.6667	0.8000	0.6667	0.6250
Medium	0.9759	0.9759	0.9759	0.9759
Medium	0.9114	0.9177	0.9114	0.9108
Medium	0.8421	0.8895	0.8421	0.8448
Medium	0.8333	0.8719	0.8333	0.8361
Long	0.9688	0.9691	0.9688	0.9688
Long	0.9242	0.9336	0.9242	0.9249
Long	0.9517	0.9545	0.9517	0.9517
Long	0.9286	0.9368	0.9286	0.9291

Table 3: Summary of Model Development Statistics

crucial for user trust. We calculate these metrics for each text slice, using the scores metrics from `sklearn.metrics`, this allows us to tailor our model's training to improve these specific metrics, depending on the needs of the deployment environment. By monitoring these metrics, we can adjust our model to ensure that it equally weighs both false positives and false negatives, which is essential for phishing detection.

**8.3.3 Robustness to Input Perturbations by Metamorphic and Differential Testing.** We further test the model's robustness by introducing slight perturbations to the text data. This simulates real-world scenarios where input data might not be perfectly formatted or might contain minor errors. The test verifies that even with these changes, the model's performance remains stable:

```
perturbed_x_train = [perturb_text(text)
for text in raw_x_train]
```

Ensuring that our model can handle such inconsistencies is critical for deploying a reliable phishing detector.

Metamorphic testing is an effective approach for this purpose, as it involves generating new test cases from existing ones by applying transformations that should not affect the expected outcomes. For example, by adding uninformative code elements or renaming variables, we can test the robustness of ML models to these variations

[1]. This approach has been shown to reveal weaknesses in models that are not evident through traditional accuracy metrics alone [1].

Additionally, in the context of phishing detection, metamorphic testing can simulate various real-world perturbations, such as typos, synonyms, or changes in text structure. This ensures that the model's performance is stable under these conditions and helps identify potential vulnerabilities that could be exploited [1, 2].

By incorporating metamorphic testing, we can enhance our model's resilience and reliability, making it better suited for deployment in environments where data inconsistencies are common [1].

**8.3.4 Model Retraining and Evaluation.** Finally, the model's ability to retrain and adapt to new data while maintaining its performance is tested. We retrain the model using different subsets of data, seeded randomly to ensure diverse training scenarios. This step is needed to ensure that the model does not overfit to a particular data configuration and can generalize well across various data distributions.

## 8.4 Tests for ML Infrastructure

Our infrastructure testing is necessary for ensuring the efficient operation of our phishing detection models, particularly in handling the intensive tasks of training and evaluating large datasets. This section details the specific tests we conduct to verify infrastructure readiness and performance.

**8.4.1 GPU Resource Availability.** We verify the availability of GPU resources, which are essential for accelerating the computation-intensive processes. This is conducted via the following test:

```
# "GPU is not available"
assert tf.config.list_physical_devices('GPU')
```

This ensures that our models run on appropriate hardware that can support their computational needs.

**8.4.2 Memory Usage Monitoring.** Memory usage is critically monitored before and after model training sessions using `psutil` to ensure it remains within acceptable limits. The relevant code snippet is:

```
process = psutil.Process(os.getpid())
# Memory in MB
memory_before = process.memory_info().rss / 1024 ** 2
```

Test Suite	Objective	Statistics
test_memory_usage	Measure memory usage before and after training to ensure it is within acceptable limits	Memory before: 691.14 MB, Memory after: 761.17 MB, Difference: 70.03 MB
test_robustness_to_noise	Evaluate the model's robustness to noisy input by comparing performance on clean and noisy validation data	Original val loss: 0.699, Noisy val loss: 0.699, Relative change: 0.00

Table 4: Summary of Test Suites, Objectives, and Statistics

```
# Model training here
memory_after = process.memory_info().rss / 1024 ** 2
assert memory_after - memory_before < 1000
```

This test helps in identifying potential memory leaks or excessive memory consumption that could hinder deployment.

**8.4.3 Model Persistence.** Testing the model's persistence involves rigorous save and load procedures to ensure that the model retains its configuration and weights across sessions, critical for deployment scenarios. We perform:

```
model.save('model.h5')
loaded_model = tf.keras.models.load_model('model.h5')
```

```
# Model mismatch on load"
assert model.get_config() == loaded_model.get_config()
```

This step verifies that the saved model when loaded, remains identical to the original, maintaining its integrity over time.

Each of these tests is integrated into our continuous integration pipeline to automatically verify system capabilities and readiness as part of our development process. This automated testing ensures that any infrastructure-related issues are identified early and can be addressed before they impact the production environment.

## 8.5 Monitoring Tests for ML

We are developing monitoring tests to maintain our model's reliability in production environments. These tests include tracking performance metrics over time, detecting anomalies in model behavior, and initiating alerts or retraining if performance drops or data drifts.

## 8.6 Suggestions for Improvement

To enhance the testing framework, a multi-faceted approach is essential. Monitoring tests need to be developed to ensure the model's long-term reliability in production environments. These tests should continuously evaluate the model's performance and integrity, alerting the team to any degradation or unexpected behaviors that may surface over time.

Additionally, the scope of robustness checks should be expanded. Incorporating stress testing with extreme values and including adversarial examples where applicable will test the model's resilience under unusual or unexpected conditions. This expansion will help identify potential weaknesses that could be exploited under real-world operating conditions, thereby strengthening the model's defenses.

Moreover, a clearer separation and more detailed isolation in tests are necessary to ensure that individual features and model

evaluations can be independently verified. This approach will allow for pinpointing specific components or processes within the model that may be underperforming or causing errors, facilitating targeted improvements and enhancing the overall robustness of the model.

## 9 CONCLUSION

- limitations
- Summary of findings and contributions
- Future work and recommendations

## REFERENCES

- [1] Leonhard Applis, Annibale Panichella, and Arie van Deursen. 2021. Assessing Robustness of ML-Based Program Analysis Tools using Metamorphic Program Transformations. In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. 1377–1381. <https://doi.org/10.1109/ASE51524.2021.9678706>
- [2] Cynthia C. S. Liem and Annibale Panichella. 2020. Run, Forest, Run? On Randomization and Reproducibility in Predictive Software Engineering. *CoRR* abs/2012.08387 (2020). arXiv:2012.08387 <https://arxiv.org/abs/2012.08387>